


ECE 5775 (Fall'17)
High-Level Digital Design Automation

Vivado HLS Tutorial

Steve Dai, Sean Lai, Zhiru Zhang
School of Electrical and Computer Engineering



Cornell University



Agenda

- ▶ Logistics and questions
- ▶ What is high-level synthesis?
- ▶ C-based synthesis
- ▶ Case study: FIR filter

High-Level Synthesis (HLS)

► What

- *Automated* design process that transforms a **high-level functional specification to optimized register-transfer level (RTL)** descriptions for efficient hardware implementation

► Why


- Productivity
 - lower design complexity and faster simulation speed
- Permutability
 - rapid design space exploration -> higher quality of result (QoR)
- Portability
 - single source -> multiple implementations

Permutability: Faster Design Space Exploration

Latency

Area


Throughput



Control-Data Flow Graph

$$out1 = f(in1, in2, in3, in4)$$

Untimed




$$t_{clk} = 3 \cdot d_{add}$$

$$T_1 = 1 / t_{clk}$$

$$A_1 = 3 \cdot A_{add}$$

Combinational




$$t_{clk} \approx d_{add} + d_{setup}$$

$$T_2 = 1 / (3 \cdot t_{clk})$$

$$A_2 = A_{add} + 2 \cdot A_{reg}$$

Sequential



$$t_{clk} = d_{add} + d_{setup}$$

$$T_3 = 1 / t_{clk}$$

$$A_3 = 3 \cdot A_{add} + 6 \cdot A_{reg}$$

Pipelined

Typical C/C++ Synthesizable Subset

- ▶ Data types:
 - Primitive types: (u)char, (u)short , (u)int, (u)long, float, double
 - Arbitrary precision integer or fixed-point types
 - Composite types: array, struct, class
 - Templated types: template<>
 - Statically determinable pointers
- ▶ No support for dynamic memory allocations
- ▶ No support for recursive function calls

Typical C/C++ Constructs to RTL Mapping


<u>C Constructs</u>		<u>HW Components</u>
Functions	→	Modules
Arguments	→	Input/output ports
Operators	→	Functional units
Scalars	→	Wires or registers
Arrays	→	Memories
Control flows	→	Control logics

Function Hierarchy


- ▶ Each function is usually translated into an RTL module
 - Functions may be inlined to dissolve their hierarchy

Source code

```
void A() { .. body A .. }  
void C() { .. body C .. }  
void B() {  
    C();  
}  
  
void TOP() {  
    A(...);  
    B(...);  
}
```




RTL hierarchy



Function Arguments


- ▶ Function arguments become ports on the RTL blocks



- Additional control ports are added to the design
- ▶ Input/output (I/O) protocols
 - Allow RTL blocks to automatically synchronize data exchange

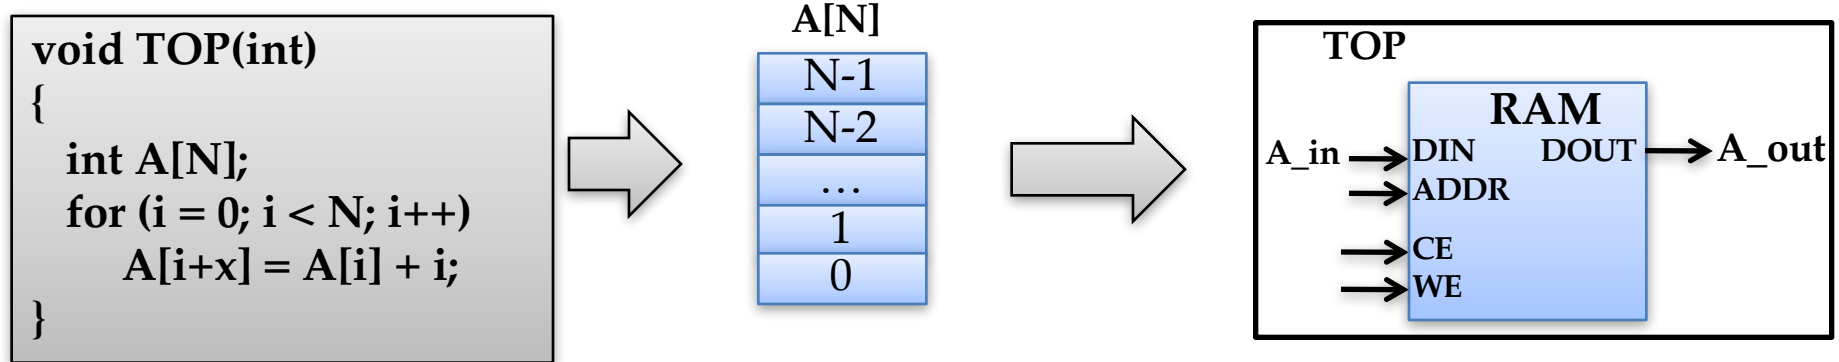
Expressions

- ▶ HLS generates datapath circuits mostly from expressions
 - Timing constraints influence the degree of registering



Arrays

- ▶ By default, an array in C code is typically implemented by a memory block in the RTL
 - Read & write array -> RAM; Constant array -> ROM





- ▶ An array can be partitioned and map to multiple RAMs
- ▶ Multiples arrays can be merged and map to one RAM
- ▶ An array can be partitioned into individual elements and map to registers

Loops

- By default, loops are rolled
 - Each loop iteration corresponds to a “sequence” of states (possibly a DAG)
 - This state sequence will be repeated multiple times based on the loop trip count

```
void TOP (...) {  
    ...  
    for (i = 0; i < N; i++)  
        b += a[i];  
}
```



Loop Unrolling

- ▶ Loop unrolling to expose higher parallelism and achieve shorter latency
 - Pros
 - Decrease loop overhead
 - Increase parallelism for scheduling
 - Facilitate constant propagation and array-to-scalar promotion
 - Cons
 - Increase operation count, which may negatively impact area, power, and timing


```
for (int i = 0; i < N; i++)  
    A[i] = C[i] + D[i];
```




```
A[0] = C[0] + D[0];  
A[1] = C[1] + D[1];  
A[2] = C[2] + D[2];  
.....
```

Loop Pipelining

- ▶ Loop pipelining is one of the most important optimizations for high-level synthesis
 - Allows a new iteration to begin processing before the previous iteration is complete
 - Key metric: **Initiation Interval (II)** in # cycles



ld – Load
st – Store



Case Study:
Finite Impulse Response (FIR) Filter

Finite Impulse Response (FIR) Filter

$$y[n] = \sum_{i=0}^N b_i x[n - i]$$

$x[n]$ input signal

$y[n]$ output signal

N filter order

b_i i th filter coefficient

```
// original, non-optimized version of FIR

#define SIZE 128
#define N 10

void fir(int input[SIZE], int output[SIZE]) {

    // FIR coefficients
    int coeff[N] = {13, -2, 9, 11, 26, 18, 95, -43, 6, 74};

    // exact translation from FIR formula above
    for (int n = 0; n < SIZE; n++) {
        int acc = 0;
        for (int i = 0; i < N; i++) {
            if (n - i >= 0)
                acc += coeff[i] * input[n - i];
        }
        output[n] = acc;
    }
}
```

Server Setup

- ▶ Log into ece-linux server
 - Host name: ecelinux.ece.cornell.edu
 - User name and password: [Your NetID credentials]
- ▶ Setup tools for this class
 - Source class setup script to setup Vivado HLS

```
> source /classes/ece5775/setup-ece5775.sh
```
- ▶ Test Vivado HLS
 - Open Vivado HLS interactive environment

```
> vivado_hls -i
```
 - List the available commands

```
> help
```


Copy FIR Example to Your Home Directory

```
> cd ~  
> cp -r /classes/ece5775/FIR_tutorial/ .  
> ls
```

- ▶ Design files
 - fir.h: function prototypes
 - fir_*.c: function definitions
- ▶ Testbench files
 - fir-top.c: function used to test the design
- ▶ Synthesis configuration files
 - run.tcl: script for configuring and running Vivado HLS

Project Tcl Script

```
#=====
# run.tcl for FIR
#=====

# open the HLS project fir.prj
open_project fir.prj -reset

# set the top-level function of the design to be fir
set_top fir

# add design and testbench files
add_files fir_initial.c
add_files -tb fir-top.c

open_solution "solution1"

# use Zynq device
set_part xc7z020clg484-1

# target clock period is 10 ns
create_clock -period 10
```

```
# do a c simulation
csim_design

# synthesize the design
csynth_design

# do a co-simulation
cosim_design

# close project and quit
close_project

# exit Vivado HLS
quit
```

You can use multiple Tcl scripts to automate different runs with different configurations.

Synthesize and Simulate the Design

```
> vivado_hls -f run.tcl
```

```
Generating csim.exe  
128/128 correct values!  
INFO: [SIM 211-1] CSim done with 0 errors.
```

```
INFO: [HLS 200-10] -----  
INFO: [HLS 200-10] -- Scheduling module 'fir'  
INFO: [HLS 200-10] -----
```

```
INFO: [HLS 200-10] -----  
INFO: [HLS 200-10] -- Exploring micro-architecture for module 'fir'  
INFO: [HLS 200-10] -----
```

```
INFO: [HLS 200-10] -----  
INFO: [HLS 200-10] -- Generating RTL for module 'fir'  
INFO: [HLS 200-10] -----
```

```
INFO: [COSIM 212-47] Using XSIM for RTL simulation.  
INFO: [COSIM 212-14] Instrumenting C test bench ...
```

```
INFO: [COSIM 212-12] Generating RTL test bench ...  
INFO: [COSIM 212-323] Starting verilog simulation.  
INFO: [COSIM 212-15] Starting XSIM ...
```

```
INFO: [COSIM 212-316] Starting C post checking ...  
128/128 correct values!
```

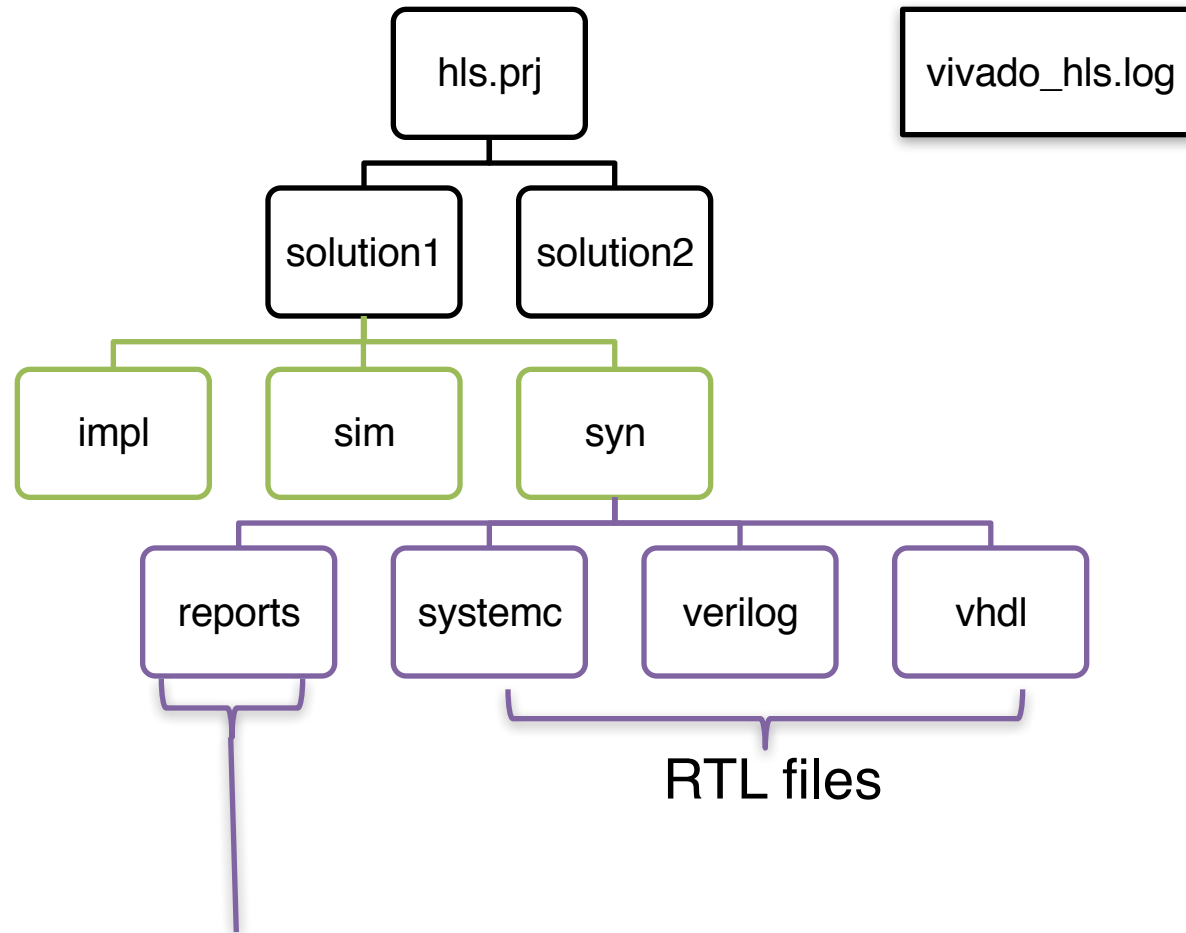
```
INFO: [COSIM 212-1000] *** C/RTL co-simulation finished: PASS ***
```

SW simulation only.
Same as simply running a
software program.

HLS
Synthesizing C to RTL

HW-SW co-simulation.
SW test bench invokes RTL
simulation.

Synthesis Directory Structure

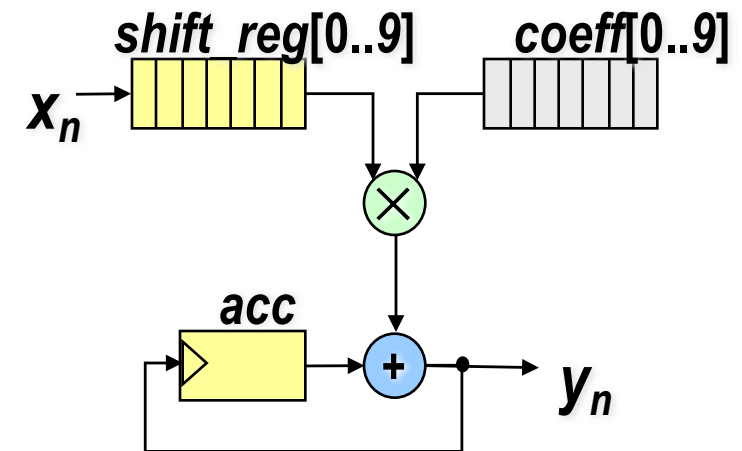


Synthesis reports of each function in the design, except those inlined.

Default Microarchitecture

```
void fir(int input[SIZE], int output[SIZE]) {  
    // FIR coefficients  
    int coeff[N] = {13, -2, 9, 11, 26, 18, 95, -43, 6, 74};  
    // Shift registers  
    int shift_reg[N] = {0, 0, 0, 0, 0, 0, 0, 0, 0, 0};  
    // loop through each output  
    for (int i = 0; i < SIZE; i++) {  
        int acc = 0;  
        // shift registers  
        for (int j = N - 1; j > 0; j--) {  
            shift_reg[j] = shift_reg[j - 1];  
        }  
        // put the new input value into the first register  
        shift_reg[0] = input[i];  
        // do multiply-accumulate operation  
        for (j = 0; j < N; j++) {  
            acc += shift_reg[j] * coeff[j];  
        }  
        output[i] = acc;  
    }  
}
```

$$y[n] = \sum_{i=0}^N b_i x[n - i]$$



Possible optimizations

- Loop unrolling
- Array partitioning
- Pipelining

Unroll Loops

```
void fir(int input[SIZE], int output[SIZE]) {
```

```
...
```

```
// loop through each output
```

```
for (int i = 0; i < SIZE; i ++ ) {  
    int acc = 0;
```

```
    // shift the registers
```

```
    for (int j = N - 1; j > 0; j--) {  
        #pragma HLS unroll  
        shift_reg[j] = shift_reg[j - 1];  
    }
```

```
...
```

```
// do multiply-accumulate operation
```

```
for (j = 0; j < N; j++) {  
    #pragma HLS unroll  
    acc += shift_reg[j] * coeff[j];  
}
```

```
...
```

```
}  
}
```

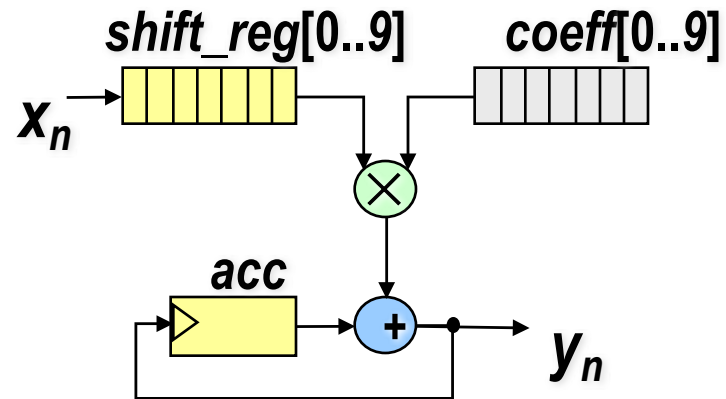
Complete
unrolling

```
// unrolled shift registers  
shift_reg[9] = shift_reg[8];  
shift_reg[8] = shift_reg[7];  
shift_reg[7] = shift_reg[6];  
...  
shift_reg[1] = shift_reg[0];
```

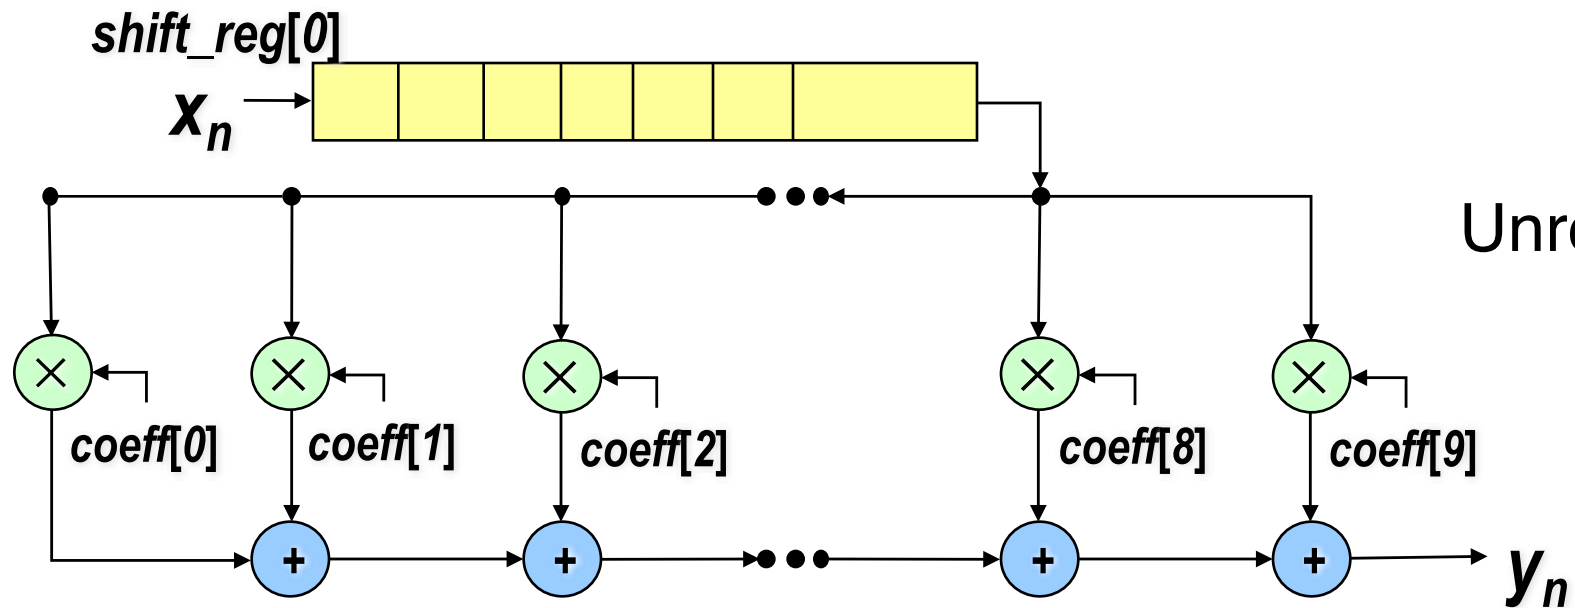
Complete
unrolling

```
// unrolled multiply-accumulate  
acc += shift_reg[0] * coeff[0];  
acc += shift_reg[1] * coeff[1];  
acc += shift_reg[2] * coeff[2];  
...  
acc += shift_reg[9] * coeff[9];
```

Microarchitecture after Unrolling



Default



Unrolled

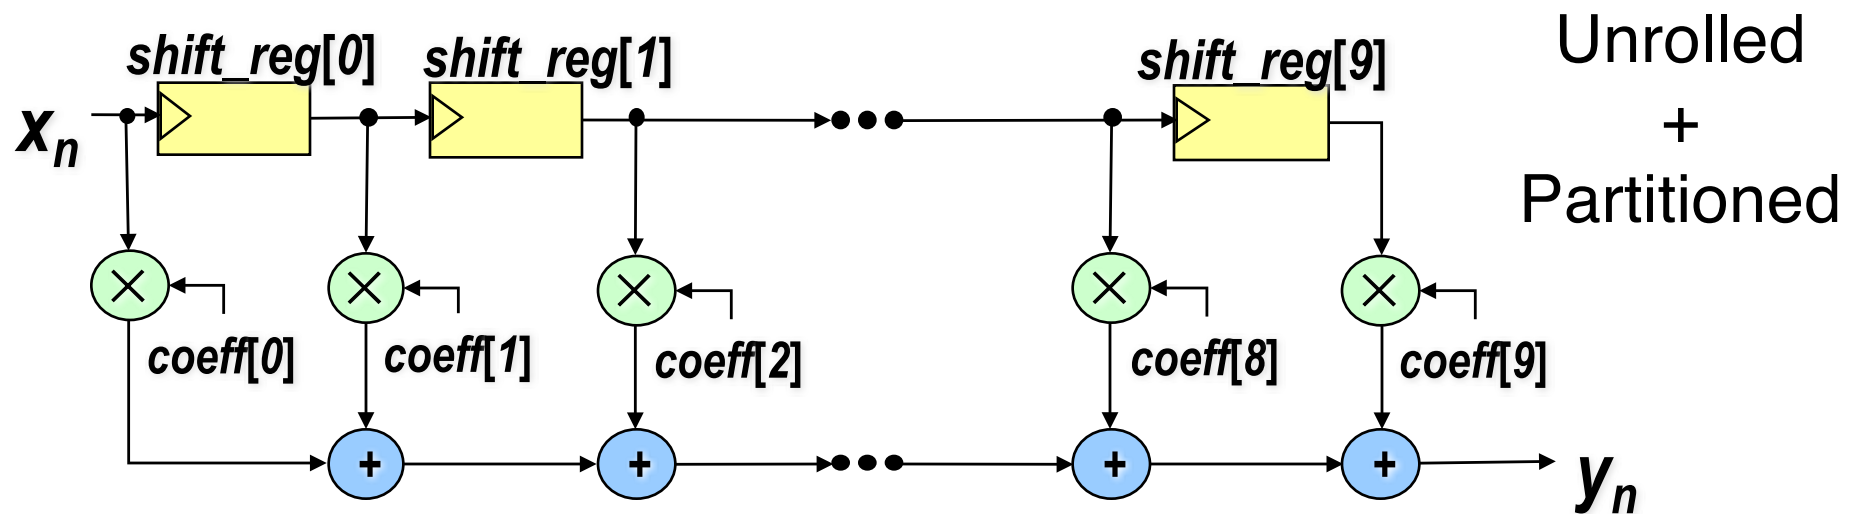
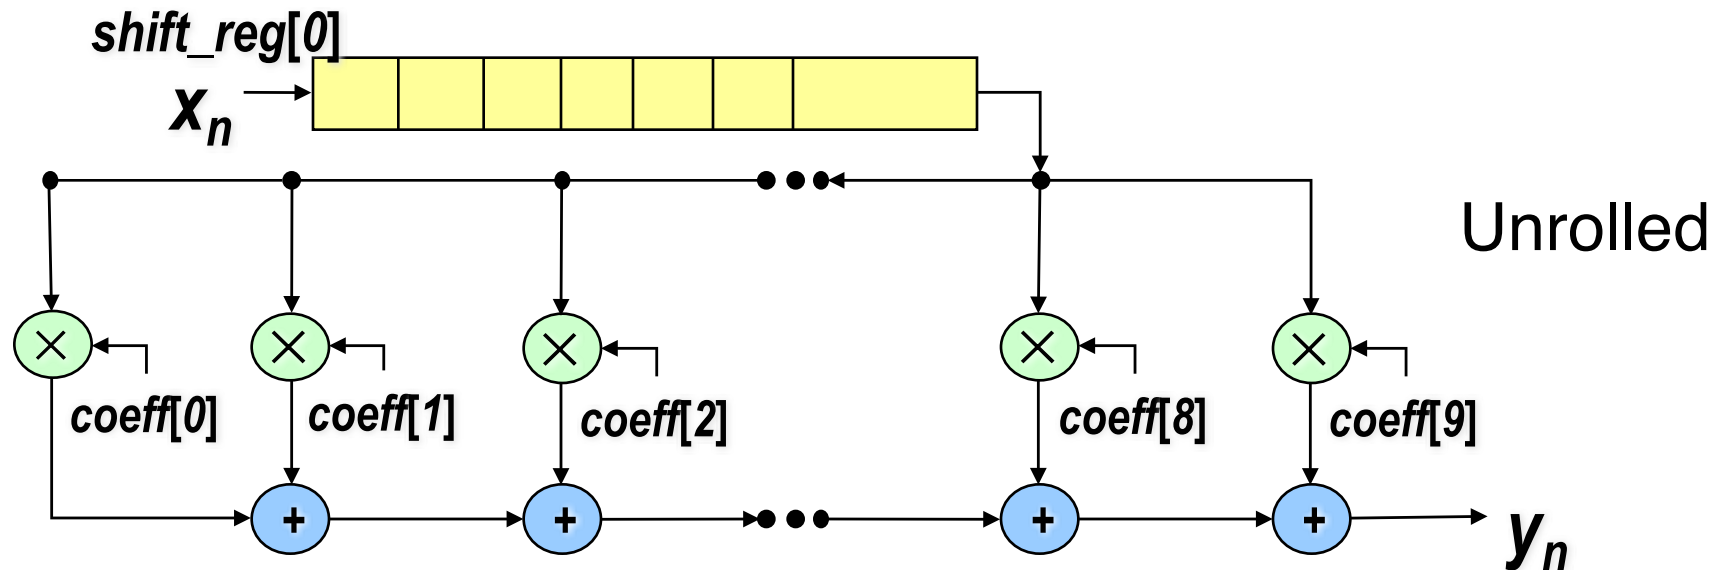
Partition Arrays

```
void fir(int input[SIZE], int output[SIZE]) {  
    // FIR coefficients  
    int coeff[N] = {13, -2, 9, 11, 26, 18, 95, -43, 6, 74};  
    // Shift registers  
    int shift_reg[N] = {0, 0, 0, 0, 0, 0, 0, 0, 0, 0};  
    #pragma HLS ARRAY_PARTITION variable=shift_reg complete dim=0  
    ...  
}
```

Complete array partitioning

```
// Shift registers  
int shift_reg_0 = 0;  
int shift_reg_1 = 0;  
int shift_reg_2 = 0;  
...  
int shift_reg_9 = 0;
```


Microarchitecture after Partitioning



Pipeline Outer Loop

```
void fir(int input[SIZE], int output[SIZE]) {
```

```
...
```

```
// loop through each output
```

```
for (int i = 0; i < SIZE; i++) {
```

```
#pragma HLS pipeline II=1
```

```
int acc = 0;
```

```
// shift the registers
```

```
for (int j = N - 1; j > 0; j--) {
```

```
#pragma HLS unroll
```

```
shift_reg[j] = shift_reg[j - 1];
```

```
}
```

```
...
```

```
// do multiply-accumulate operation
```

```
for (j = 0; j < N; j++) {
```

```
#pragma HLS unroll
```

```
acc += shift_reg[j] * coeff[j];
```

```
}
```

```
...
```

```
}
```

```
}
```

Pipeline the entire outer loop

```
// loop through each output
```

```
for (int i = 0; i < SIZE; i++) {
```

```
#pragma HLS pipeline II=1
```

```
int acc = 0;
```

```
...
```

```
// put the new input value into the
```

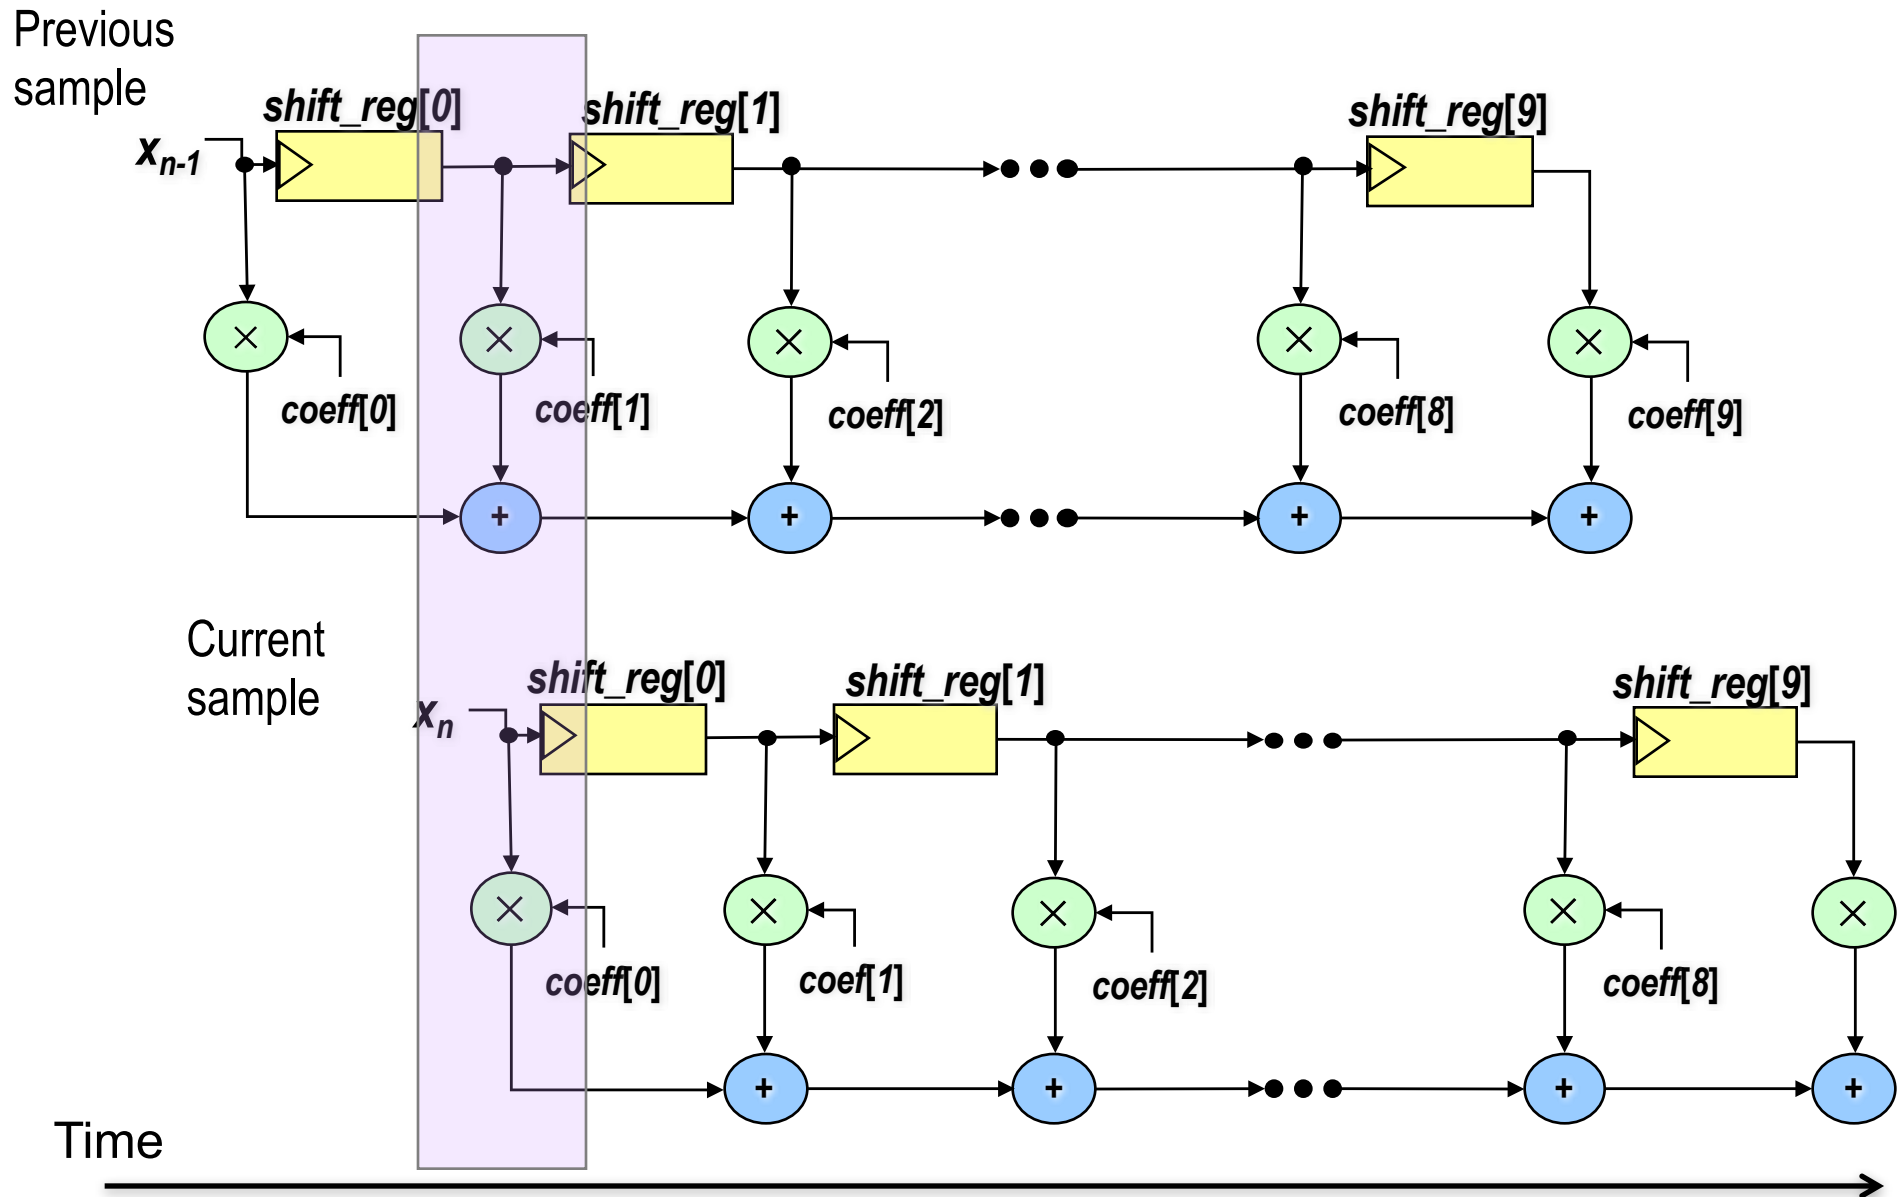
```
// first register
```

```
shift_reg[0] = input[i];
```

```
...
```

```
}
```

Fully Pipelined Implementation



Pipeline Outer Loop

```
void fir(int input[SIZE], int output[SIZE]) {  
    ...  
    // loop through each output  
    for (int i = 0; i < SIZE; i ++ ) {  
        #pragma HLS pipeline II=1  
        int acc = 0;  
        // shift the registers  
        for (int j = N - 1; j > 0; j--) {  
            #pragma HLS unroll  
            shift_reg[j] = shift_reg[j - 1];  
        }  
        ...  
        // do multiply-accumulate operation  
        for (j = 0; j < N; j++) {  
            #pragma HLS unroll  
            acc += shift_reg[j] * coeff[j];  
        }  
        ...  
    }  
}
```

Pipeline the entire outer loop

Inner loops automatically
unrolled when pipelining the
outer loop